

Meridian

Sample problem for Tribe AI

2026-10-05

TOC

- Current state and direction
- Point of view and MVP/v1 direction
- Demo

The background of the slide is a dark blue gradient. It features several horizontal, wavy lines in a lighter shade of blue, creating a layered, mountain-like or ocean-like effect. The text is centered horizontally and positioned in the upper half of the slide.

Current state and findings

Potential regulatory gap first

Up to 437 claims in the sample needed a human by regulation and may never have been routed to one.

- Two flags on every claim
 - *requires_human_by_regulation*: set by rules
 - *flagged_for_human_review*: a second-opinion signal, from a person or low AI confidence
- **437 claims (8.7%)** have the regulation flag and no review flag
 - **85** are over \$10K in one of the 12 states: the stated rule, clearly required
 - **352** are in 8 other states (AZ, IN, MO, NC, TN, TX, WA, WI), 340 of them \$10K or under. No known rule explains them.

The 12 states: 8 named in discovery (CA, NY, NJ, FL, IL, PA, OH, GA), 4 inferred from the data (MA, MD, MI, VA: every claim over \$10K is flagged, none under). Confirm the list.

Potential regulatory gap first

- **The question for engineering:** which flag decides whether a claim reaches a human?
 - **Either flag:** nothing is skipped, but the 352 reach humans for reasons nobody can explain. That's a feasibility problem for the review-rate target, not a compliance one.
 - **Review flag only:** at least 85 required reviews skipped, ~6,800 a year at 400K volume; up to 437, ~35,000 a year.
 - **Neither directly:** Something other than these flags routes to human agents.
- **The fix, if needed:** route to a human when *either* flag is set. One condition in the routing rule. Value before any AI.
- Depending upon how the <20% SOW objective is defined, it might be infeasible.

requires_human_by_regulation

flagged_for_human_review

	TRUE	FALSE
TRUE	835 (16.7%) Required human review	437 (8.7%) 85 > \$10k in 12 states 352 in other states, unexplained
FALSE	1,445 (28.9%) Second opinion only	2,283 (45.7%) No human needed

Today (review flag only) => $835 + 1,445 = 2,280$ (45.6%)

Either flag, stated rule => $2,280 + 85$ (47.3%)

Either flag, labels as they stand => $2,280 + 437$ (85 + 352) (54.3%)

IF we only consider those that require a second opinion review when looking at the <20% number, then about 40-45% of the total claim volume can be expected to require human review instead of the 54.3% in the current sample and rules. Still a meaningful drop in volume across 400k claims a year, but expectations need to be aligned across parties.

Tension to name before Pooja does: the fix raises the review rate (45.6% → 47.3% under the stated rule; 54.3% on the labels as they stand).

The 12 states: 8 named in discovery (CA, NY, NJ, FL, IL, PA, OH, GA), 4 inferred from the data (MA, MD, MI, VA: every claim over \$10K is flagged, none under). Confirm the list.

The SOW targets

One target is potentially possible, the others not so much.

- Cycle time 68h → <24h → Possible, but very difficult without substantially increasing average throughput per agent
- Cost to serve → -25% blended → Near impossible. Moving *all* fax/EDI to the portal saves ≤8.8%. Moving *all* fax/EDI *and* phone to portal saves ~21.3%.
 - Unless we can account for substantial cost savings by reducing the second opinion reviews, a 25% reduction is impossible.
- Review rate 45% → <20% → Impossible or needs a new definition. NOTE: the current real total review rate is 54.3% in the sample data, suggesting there is a real gap in regulated review requirements.
 - The floor due to current regulation rules is 25.4% of total in the sample
 - We *could* reduce the second opinion human reviews from 28.9% to <20%, but that leaves ~45% of total volume human reviewed including regulation required.
- CSAT 3.2 → > 4.0 → Near impossible, likely needs to be redefined
 - If denials stay at 2.24, the other 2,421 respondents would need to average 4.39 for the overall figure to pass 4.0

Today, even paid-in-full claims, the best case, average 3.65, and only 113 of 2,421 non-denied responses (4.7%) are a 5. An average of 4.39 means most claimants giving a 5, and today almost none do. Even if every non-denied claimant gave a 5, the overall average would top out around 4.5.

41% of claims have no score. The 3.2 baseline comes only from people who chose to respond (reference/data_dictionary.md:133). If the response rate changes, the number can move with no change in service. For example, a new survey channel might bring in more unhappy respondents, or more happy ones.

Speed alone won't carry the satisfaction.



MVP and beyond

Point of view

Make the car run better; don't change the dashboard

- **Who:** 95 adjusters on ClaimsPro, eight screens deep for a claim. The primary persona is the second-year adjuster.
- **The point of view:** fix what happens *before* the adjuster opens a claim. Focus on intake, validation, prioritization, routing to the right person.
- **The human stays in the loop:** the adjuster makes every decision (deny in particular).
- **The ideal:** an adjuster opens ClaimsPro and sees one thing, the next right claim for them.

- For example, a single claim took 208 hours. 117 claims in a queue.
 - Show the result in ClaimsPro's own fields. Add new interface only where the existing one hits a wall.
 - The system prepares, orders, explains, and records; it never approves or denies.
 - ClaimsPro can't be made to work that way, but every step moves toward it.
-
- Ground it in 4222 before any numbers: phone intake, a rough auto-transcript, a timeline built by hand, unsure on coverage, four days in senior review (Ops, 56:10–62:15).
 - "Don't build us an eighth screen" (Ops, 86:20). Sandra's own bet: "half my cycle-time problem is a dispatch problem" (Ops, 20:15).
 - Sources: `sample\_claims/IS-CLM-2025004222/`; `memory/calltranscripts/w1-tue-sandra-ops-walkthrough.md`.

Five ways to work with a legacy system

- Replace ClaimsPro entirely
- Replace the UI; ClaimsPro as a headless backend
- Build inside ClaimsPro with a vendor SDK or plugin
- Write into ClaimsPro's own fields (notes, custom fields)
- Panel beside ClaimsPro (managed browser extension)

- This is strategy, not plumbing: it rules out every design that depends on replacing or re-skinning ClaimsPro. The next slide is the design choices made inside this box.
- 4 and 5 work together, in that order: custom fields come first, as the record and the adjuster's v1 view (they also feed ClaimsPro's own rules and reports). The overlay comes only when the fields hit a wall. If the extension is ever down, the fields still show in ClaimsPro.
- Also rejected: bots driving the ClaimsPro screens. They break on every vendor release and would route around the deliberate decisioning gap.
- Not a UI path, but worth one line: fixing fax OCR *before* it enters ClaimsPro (it lands today with no human check).
- If OCR correction comes up: today the LLM may fix characters only, never words, and code checks the corrected text word by word against the raw OCR. On 2993, dropping `n0t` as noise would flip "not consistent with a fresh impact", the evidence behind the denial. Horizon: models that predict meaning rather than the next token (JEPA-style) may handle this more reliably; none are ready for documents, so it's one to watch over the next couple of years, not to build on now.
- "Don't build us an eighth screen" (Ops, 86:20). Neither path adds one: the fields live in ClaimsPro, and the overlay sits on the screen adjusters already use.
- Sources: Systems, 4:40 and 5:45; clarification call (SaaS, no SDK, custom fields, force-installable extensions).

Five ways to work with a legacy system

- **Write into ClaimsPro's own fields** (notes, custom fields)
 - Custom fields exist; SOAP writes cover notes and documents. Durable, auditable, survives if anything else breaks
- **Panel beside ClaimsPro** (managed browser extension)
 - SaaS in a browser; IT can force-install. Reads the open claim, shows our panel, calls our backend

- This is strategy, not plumbing: it rules out every design that depends on replacing or re-skinning ClaimsPro. The next slide is the design choices made inside this box.
- 4 and 5 work together, in that order: custom fields come first, as the record and the adjuster's v1 view (they also feed ClaimsPro's own rules and reports). The overlay comes only when the fields hit a wall. If the extension is ever down, the fields still show in ClaimsPro.
- Also rejected: bots driving the ClaimsPro screens. They break on every vendor release and would route around the deliberate decisioning gap.
- Not a UI path, but worth one line: fixing fax OCR *before* it enters ClaimsPro (it lands today with no human check).
- If OCR correction comes up: today the LLM may fix characters only, never words, and code checks the corrected text word by word against the raw OCR. On 2993, dropping `n0t` as noise would flip "not consistent with a fresh impact", the evidence behind the denial. Horizon: models that predict meaning rather than the next token (JEPA-style) may handle this more reliably; none are ready for documents, so it's one to watch over the next couple of years, not to build on now.
- "Don't build us an eighth screen" (Ops, 86:20). Neither path adds one: the fields live in ClaimsPro, and the overlay sits on the screen adjusters already use.
- Sources: Systems, 4:40 and 5:45; clarification call (SaaS, no SDK, custom fields, force-installable extensions).

What we're not building yet, and when it would come back

Rejected for now	Why	Where it could come back
A separate app in another tab	Adds the ninth screen; splits attention	Managers only: the pipeline monitor is its own page. A link from ClaimsPro's own menus is P3, only if adoption lags.
Rebuilding the eight-screen experience	No decision API, SOAP-only writes; not doable in six months	A side-panel summary and table of contents: reads the claim, raises what matters, links to where each point lives.
A chat agent	An interface looking for a problem; we try no new interface first	P2/P3, in the same side panel. Also the visible "AI" executives can point to.
A training and precedent tool	Pre-processing notes in custom fields cover part of the need first	Added to the side panel later: the second-year adjuster's top ask

Sequence: no new adjuster UI (pre-processing, routing, custom fields) → one side panel (summary and table of contents) → the same panel gains research, then chat if it earns a place.

No new interface for adjusters until the existing one hits a wall.

The assumption everything rests on: ClaimsPro's custom fields are prominent on the screens adjusters already use, and make sense there. The first version is built on that. If it's wrong, or the fields read badly in the legacy UI, the same content moves into the side panel for display, and stays stored as custom fields in ClaimsPro.

- Name the tension before Carlos does: "We spent \$5M and nobody sees anything different." The answer: the change is visible where it's measured, in the pipeline monitor for Sandra, leadership, and us. The side panel is the tangible thing on the adjuster's screen, once routing has earned trust.
- Also rejected: auto-approval. A one-click, human-confirmed fast lane instead. Consent for automated adjudication is untracked in 8 states, and claims over \$10K need human review in 12.
- Precedent search is deferred, not dismissed. Built from old notes, it copies past judgment forward, bias included; the NAIC bias-testing commitment applies.
- Open question for Michael: where do custom fields render (I3 confirmed they exist, not where they show)?

Who this changes, and how we'll know

Five groups, three flows that change, and a measure for each at design time and in production.

Who	What they need	What changes for them
Adjusters	Claims that arrive ready to decide; a reference for hard calls; no eighth screen	Queue arrives ordered; custom fields say why a claim is theirs and what's prepared or missing
Senior adjusters and leads	Time back for complex files; fewer interruptions	Hard claims arrive already routed to them; fewer "can you look at this" escalations
Sandra, claims ops	Cycle time down without burning people out	A dispatcher she can tune; visibility into why claims sit
Michael and Compliance	"Show me how you know it's right, and keep showing me"	Every route and summary explained and logged; the regulatory gap closed; drift visible
Carlos	Cost down, demonstrated not promised	A measured path on the targets that are reachable

Not direct users but affected: claimants (faster, better-explained outcomes) and the call center (phone intake feeds the same brief later).

Flows

Flow	Change	Design-time measure	Production measure (OKR)
Routing: claim arrives → queue → adjuster	<i>Change:</i> round-robin becomes ordered and matched; routes on both flags	Replay the extract: simulated queue wait vs actual; 0 regulated claims unrouted; adjusters agree with the route on sampled claims	O: Claims reach the right adjuster without waiting. KR: simple-claim median cycle <24h (today 34.7h; 34% under 24h); 0 regulated claims without a review step; escalations to senior review down X%
Preparation: claim opened → ready to decide	<i>Create:</i> pre-processing results in custom fields (v1); the side-panel summary and table of contents (v2)	Adjuster task test on sample claims: time to first decision, fields re-keyed, "I trust this" rating; extraction accuracy vs adjuster-corrected values	O: Adjusters decide, not reassemble. KR: handling minutes per claim down Y%; second opinions down Z%; brief accuracy ≥ agreed bar on audited sample
Correction: adjuster disagrees → system learns	<i>Create:</i> the learning loop	Corrections take one action; a rule fix finds the similar claims in bulk (the 85)	O: The system gets more right each month, visibly. KR: correction rate trending down; drift alerts reviewed within a week; no silent rule changes; fax OCR errors reaching ClaimsPro down N%

- X% targets are placeholders, agreed in the scope document with Sandra and Michael; the baselines are real and from the extract.
- Change management (Nikolina): pilot with one team and the adjusters who shaped it; nothing takes a decision away from them; the old score comes off the screen only once the new routing has earned trust (M6).
- Every production measure needs a baseline agreed before launch, or month 3 becomes an argument about numbers.
- OCR KR: the ~78% accuracy baseline is Meridian's own figure, unverified. Measuring it needs OCR output linked to the source images and adjuster corrections captured (D3). Each adjuster fix to an extracted field is a correction the loop learns from; the 2022 OCR engine itself stays as is.

Where the first pass goes

Get the right claim to the right adjuster in the right order, ready to decide

- **Pre-processing and routing (deep):** validate intake and catch what's missing; route on both flags; order queues by age and deadline; send hard claims to the right adjuster first. Moves cycle time.
- **Custom fields (the adjuster's view, v1):** why the claim was routed, its review lane, what's been prepared or is missing. The record in ClaimsPro, and all the adjuster sees change.
- **Pipeline monitor (managers and us):** claims moving through intake, validation, enrichment, prioritization, assignment; exceptions that need a human before an adjuster sees them.


- Why routing: Sandra's own first fix (Ops, 20:15), and it moves the one target within reach.
- Human in the loop: the adjuster owns every decision. Exceptions go to a person, never to a default.
- Learning loop example: fix the routing rule once, and the 85 are caught in bulk.

How we got here

Sources of research and conclusions

- **Read the record:** SOW, three discovery calls, the analytics memo, the three sample claims.
- **Wrote down every question**, with who could answer it and why it mattered; a clarification call closed the ones that blocked design (docs/discovery/open_questions.md).
- **Tested the claims in the data**, in a script anyone can re-run, instead of trusting the memo or ourselves (analysis/review_floor.py). This is where the flag gap and the target reads came from.
- **Took positions, then revised them:** auto-approval became a one-click fast lane; skill routing moved from P3 to P2; the review floor changed from 25% to 13% once the stated rule came in.
- **Picked one adjuster and one claim** and built it end to end, rather than sketching everything.
- GitHub repo with history, scripts validating conclusions, and additional detail available at <https://github.com/joshjhall/meridian>

- For Charles: the iterations matter more than the steps. Have the before/after of each revised position ready.
- Odd patterns found along the way: the flag gap; 628 regulated labels with no rule; 20 of the 500 calibration labels still pending; `handling\_hours` is time open, not effort (reference/data_dictionary.md:93).
- Sources: git history of this repo; `docs/discovery/open\_questions.md`.



Demo